



HIGH PERFORMANCE MACHINE LEARNING

Caspar van Leeuwen
HPC consultant
SURFsara

SURF

Program, 2nd day

- 9:00-9:45 Parallel computing for Deep Learning
- 9:45-10:30 Hardware & Hardware bottlenecks
- 10:30-10:45 Coffee break
- 10:45-11:30 Profiling: creation and interpretation (Valeriu Codreanu)
- 11:30 – 12:00 Frameworks
- 12:00-13:00 Lunch
- 13:00-13:30 Horovod practical (Damian Podareanu)
- 13:30-14:30 CNN data distributed practical with Cifar 10 (Valeriu Codreanu / Damian Podareanu)
- 14:30-14:45 Coffee break
- 14:45-15:15 CNN data distributed practical with CIFAR 10 (Valeriu Codreanu)
- 15:15-15:30 Hybrid parallelism (Sagar Dolas)
- 15:30-16:15 Mesh Tensorflow tutorial (Sagar Dolas)
- 16:16-... Questions, Open Discussion, etc

Hardware

Goals:

- Understand what hardware bottlenecks could be limiting
- Understand pro's and con's of various hardware
- Know how to choose appropriate hardware for you DL task
- Know what to do to mitigate bottlenecks

Hardware bottlenecks

- Compute (floating point operations per second, FLOPS)
- Memory bandwidth
- Memory size
- I/O
- Communication

Hardware bottlenecks

- Compute (floating point operations per second, FLOPS)
- Memory bandwidth
- Memory size
- I/O
- Communication

E.g. training a compute intensive network on a single node

Hardware bottlenecks

- Compute (floating point operations per second, FLOPS)
 - Memory bandwidth
 - Memory size
 - I/O
 - Communication
- Data needs to get to the processor in time in order to do compute!
 - Many codes are limited by memory bandwidth

Hardware bottlenecks

- Compute (floating point operations per second, FLOPS)
 - Memory bandwidth
 - Memory size
 - I/O
 - Communication
- Very deep or wide networks, or networks with very large input/output layers (e.g. high resolution images) may be limited by memory size.
 - Not a performance bottleneck, but a no-go!

Hardware bottlenecks

- Compute (floating point operations per second, FLOPS)
 - Memory bandwidth
 - Memory size
 - I/O
 - Communication
- HPC systems typically have shared file systems, usually with good bandwidth, but (relatively) low IOPS
 - (Very) common bottleneck in distributed learning! Many nodes reading from the same filesystem.
 - Other users (& sysadmins) will dislike you if you do I/O in a naive way!

Hardware bottlenecks

- Compute (floating point operations per second, FLOPS)
- Memory bandwidth
- Memory size
- I/O
- Communication


Communication can be limiting in several ways:

- Latency (many, small message send between nodes)
- Bandwidth (few, large messages send between nodes)
- Load imbalance (some workers in distributed job are slower / have more work; others have to wait when synchronization is needed)

Hardware overview

A look at the hardware, from a DL perspective:

- Nvidia Pascal / Volta GPUs
- AMD Vega / Vega20
- Intel Xeon Scalable
- AMD Zen
- Specialized hardware
- I/O
- Interconnects



Source: Demystifying Parallel and Distributed Deep Learning: An In-Depth Concurrency Analysis, Ben-Nun & Hoefler 2018

Hardware overview

	INT8 [TOPS]	FP16 [TFLOPS]	FP32 [TFLOPS]	FP64 [TFLOPS]	Memory [GB]	Memory Bandwidth [GB/s]	PCIe [GB/s]	Proprietary Interconnect [GB/s]
AMD MI25	-	24.6	12.29	0.77	16	484	15.75	-
AMD MI50	53.6	26.8	13.4	6.7	16	1024	31.51	200 (2 × 100)
AMD MI60	58.9	29.5	14.7	7.4	32	1024	31.51	200 (2 × 100)
NVIDIA P100	-	21.2	10.6	5.3	16	732	15.75	160 (4 × 40)
NVIDIA V100	62.8	31.4 / 125	15.7	7.8	16/32	900	15.75	300 (6 × 50)
Intel Xeon Scalable 8180 (per socket)	-	-	3.0 / 4.2	1.5 / 2.1	768 (max)	119 (max)	15.75	-
AMD EPYC 7601 (per socket)	-	-	1.1 / 1.4	0.56 / 0.69	2000 (max)	159 (max)	15.75	-

Source: Prace best practice guide – deep learning

<http://www.prace-ri.eu/IMG/pdf/Best-Practice-Guide-Deep-Learning.pdf>

Hardware overview

GPUs support reduced precision, has higher performance

	INT8 [TOPS]	FP16 [TFLOPS]	FP32 [TFLOPS]	FP64 [TFLOPS]	Memory [GB]	Memory Bandwidth [GB/s]	PCIe [GB/s]	Proprietary Interconnect [GB/s]
AMD MI25	-	24.6	12.29	0.77	16	484	15.75	-
AMD MI50	53.6	26.8	13.4	6.7	16	1024	31.51	200 (2 × 100)
AMD MI60	58.9	29.5	14.7	7.4	32	1024	31.51	200 (2 × 100)
NVIDIA P100	-	21.2	10.6	5.3	16	732	15.75	160 (4 × 40)
NVIDIA V100	62.8	31.4 / 125	15.7	7.8	16/32	900	15.75	300 (6 × 50)
Intel Xeon Scalable 8180 (per socket)	-	-	3.0 / 4.2	1.5 / 2.1	768 (max)	119 (max)	15.75	-
AMD EPYC 7601 (per socket)	-	-	1.1 / 1.4	0.56 / 0.69	2000 (max)	159 (max)	15.75	-

Source: Prace best practice guide – deep learning

<http://www.prace-ri.eu/IMG/pdf/Best-Practice-Guide-Deep-Learning.pdf>

Hardware overview

GPUs have a lot of FLOPS compared to CPUs

	INT8 [TOPS]	FP16 [TFLOPS]	FP32 [TFLOPS]	FP64 [TFLOPS]	Memory [GB]	Memory Bandwidth [GB/s]	PCIe [GB/s]	Proprietary Interconnect [GB/s]
AMD MI25	-	24.6	12.29	0.77	16	484	15.75	-
AMD MI50	53.6	26.8	13.4	6.7	16	1024	31.51	200 (2 × 100)
AMD MI60	58.9	29.5	14.7	7.4	32	1024	31.51	200 (2 × 100)
NVIDIA P100	-	21.2	10.6	5.3	16	732	15.75	160 (4 × 40)
NVIDIA V100	62.8	31.4 / 125	15.7	7.8	16/32	900	15.75	300 (6 × 50)
Intel Xeon Scalable 8180 (per socket)	-	-	3.0 / 4.2	1.5 / 2.1	768 (max)	119 (max)	15.75	-
AMD EPYC 7601 (per socket)	-	-	1.1 / 1.4	0.56 / 0.69	2000 (max)	159 (max)	15.75	-

Source: Prace best practice guide – deep learning

<http://www.prace-ri.eu/IMG/pdf/Best-Practice-Guide-Deep-Learning.pdf>

Hardware overview

CPU's have a lot of memory compared to GPU's

	INT8 [TOPS]	FP16 [TFLOPS]	FP32 [TFLOPS]	FP64 [TFLOPS]	Memory [GB]	Memory Bandwidth [GB/s]	PCIe [GB/s]	Proprietary Interconnect [GB/s]
AMD MI25	-	24.6	12.29	0.77	16	484	15.75	-
AMD MI50	53.6	26.8	13.4	6.7	16	1024	31.51	200 (2 × 100)
AMD MI60	58.9	29.5	14.7	7.4	32	1024	31.51	200 (2 × 100)
NVIDIA P100	-	21.2	10.6	5.3	16	732	15.75	160 (4 × 40)
NVIDIA V100	62.8	31.4 / 125	15.7	7.8	16/32	900	15.75	300 (6 × 50)
Intel Xeon Scalable 8180 (per socket)	-	-	3.0 / 4.2	1.5 / 2.1	768 (max)	119 (max)	15.75	-
AMD EPYC 7601 (per socket)	-	-	1.1 / 1.4	0.56 / 0.69	2000 (max)	159 (max)	15.75	-

Source: Prace best practice guide – deep learning

<http://www.prace-ri.eu/IMG/pdf/Best-Practice-Guide-Deep-Learning.pdf>

Hardware overview

Memory bandwidth relative to FLOPS is approximately the same

	INT8 [TOPS]	FP16 [TFLOPS]	FP32 [TFLOPS]	FP64 [TFLOPS]	Memory [GB]	Memory Bandwidth [GB/s]	PCIe [GB/s]	Proprietary Interconnect [GB/s]
AMD MI25	-	24.6	12.29	0.77	16	484	15.75	-
AMD MI50	53.6	26.8	13.4	6.7	16	1024	31.51	200 (2 × 100)
AMD MI60	58.9	29.5	14.7	7.4	32	1024	31.51	200 (2 × 100)
NVIDIA P100	-	21.2	10.6	5.3	16	732	15.75	160 (4 × 40)
NVIDIA V100	62.8	31.4 / 125	15.7	7.8	16/32	900	15.75	300 (6 × 50)
Intel Xeon Scalable 8180 (per socket)	-	-	3.0 / 4.2	1.5 / 2.1	768 (max)	119 (max)	15.75	-
AMD EPYC 7601 (per socket)	-	-	1.1 / 1.4	0.56 / 0.69	2000 (max)	159 (max)	15.75	-

Source: Prace best practice guide – deep learning

<http://www.prace-ri.eu/IMG/pdf/Best-Practice-Guide-Deep-Learning.pdf>

Nvidia Pascal / Volta GPUs

Generally *you* are responsible for specifying a reduced precision: DL frameworks don't do this automatically since it may impact your networks accuracy, convergence, etc

Features

- INT8 & FP16 support
- (Volta) Tensor cores: fused multiply-add units that support mixed precision (multiply in FP16, add in FP32). High performance: 120 TOPS.

$$D = \begin{matrix} \text{FP16 or FP32} & \begin{matrix} \text{FP16} & \text{FP16} \end{matrix} & + & \begin{matrix} \text{FP16 or FP32} \end{matrix} \end{matrix}$$

The diagram illustrates a matrix multiplication operation $D = A * B + C$. Matrix A is a 4x4 grid of elements $A_{0,0}$ through $A_{3,3}$, labeled "FP16 or FP32" below. Matrix B is a 4x4 grid of elements $B_{0,0}$ through $B_{3,3}$, labeled "FP16" below. Matrix C is a 4x4 grid of elements $C_{0,0}$ through $C_{3,3}$, labeled "FP16 or FP32" below. The result D is a 4x4 grid. The operation is shown as $D = A * B + C$.

- High bandwidth memory, 2nd generation (HBM2)
- NVLink: allows direct GPU₁ \Leftrightarrow GPU₂ communication in a multi-GPU node (much faster than having to go over PCIe)

Nvidia Pascal / Volta GPUs

(Low-level) library support

- CUDA
- cuBLAS: basic linear algebra
- cuSPARSE: sparse matrix algebra
- cuDNN: primitives for deep neural networks
- NCCL: NVIDIA collective communications library (implements efficient allreduce and supports e.g. NVLink & RDMA)

Efficient low-level libraries are just as important as hardware itself! Otherwise, theoretical hardware specs might be great, but you'll never get that performance!

AMD Vega / Vega20

Features

- INT8 & FP16 support
- High bandwidth memory, 2nd generation (HBM2)
- PCIe 4.0 support (large bandwidth CPU \Leftrightarrow GPU)
- Infinity fabric: proprietary interconnect for high bandwidth GPU \Leftrightarrow GPU communication in a multi-GPU node.

AMD Vega / Vega20

(Low-level) library support

- Heterogeneous-computing Interface for Portability (HIP): a C++ dialect that was designed to ease conversion of CUDA applications to portable C++ code
- MIOpen: machine learning primitives (based on the OpenCL or HIP)
- rocBLAS: basic linear algebra
- ROCm: software ecosystem for GPU computing
- ROCm: has forks of TensorFlow, Caffe 2, PyTorch, MxNet and CNTK with MIOpen support.

While software stack is 'young' compared to NVidia, performance is competitive – provided the ROCm forks of the frameworks are used when computing on AMD hardware!

Intel Xeon Scalable

- Supports AVX-512 vector instructions
- Supports very large memory (more than 1 TB per node)
- Intel Math Kernel Library (MKL): optimized BLAS, LAPACK and FFTW routines
- Intel MKL-DNN: primitives for deep learning.

Pro tip:

Low-end Xeon Scalable processors have only 1 AVX-512 FMA unit. AVX2 may perform better on these processors, because AVX-512 instructions are executed at lower clock speeds!

AMD Zen

- Supports AVX-2 vector instructions
- Supports very large memory (more than 1 TB per node)
- BLIS: a BLAS implementation optimized for AMD EPYC processors
- libFLAME: portable library for dense matrix multiplications

Note: the Zen processor is not really marketed as a processor for DL task. DL frameworks don't generally support its low level libraries.

More info: PRACE AMD EPYC best practice guide, <http://www.prace-ri.eu/IMG/pdf/Best-Practice-Guide-AMD.pdf>

Specialized hardware

Tensor processing units (TPU)

- Chip designed specifically for machine learning (tensor operations)
- Available only in Google Cloud
- Google Cloud TPU v3: 420 TFLOPS, 128 GB HBM. About 8 USD/TPU-hour
- Supported in TensorFlow

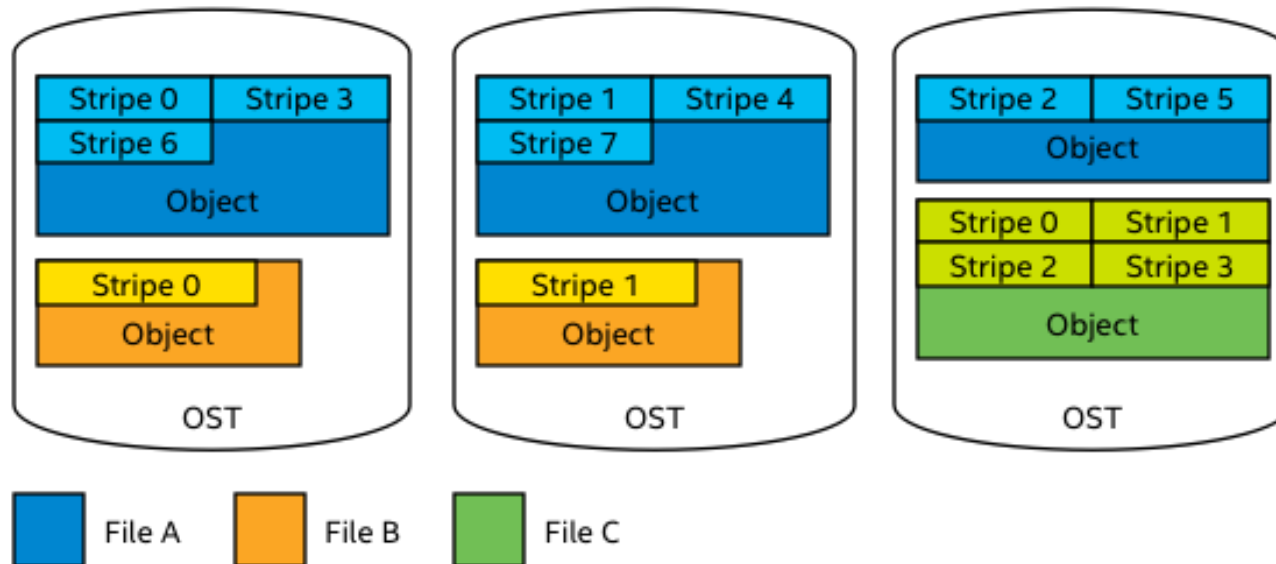
Field Programmable Gate Arrays (FPGA)

- Programmable chips
- Can do e.g. compute in any desired accuracy
- Very experimental, no framework support
- Maybe in the future...

I/O

HPC systems typically use parallel shared file systems.

- Parallel file system: one file can be distributed over many physical disks, to increase I/O bandwidth.



I/O

Two main types of parallel filesystems

- Lustre
 - Metadata (filename, size, location, etc) stored on separate server
 - Object Storage Target (OST) stores actual file
 - Striping over multiple OSTs can be managed by user
- GPFS
 - Metadata and actual file stored on the same server
 - Striping is managed automatically, by file system. User has no control.

Mitigate hardware bottlenecks

- Compute (floating point operations per second, FLOPS)
- Memory bandwidth
- Memory size
- I/O
- Communication

- Use high-FLOP hardware (GPUs, TPUs)
- Use specialized vector instructions: AVX, AVX2, AVX-512 (CPUs)
- Use reduced precision training (GPUs)
- Use distributed learning

Pro tip for Tensorflow CPU training

‘Pip install tensorflow’ installs a TensorFlow binary with only AVX support. Poor performance on new CPU architectures!

- Build from source yourself (tricky, but allows perfect optimization for your system)
- Use prebuilt Intel optimized Tensorflow on Intel CPUs: pip install intel-tensorflow, or conda. (easy, but TF versions may lag behind) See <https://software.intel.com/en-us/articles/intel-optimization-for-tensorflow-installation-guide>

Mitigate hardware bottlenecks

- Compute (floating point operations per second, FLOPS)
 - Memory bandwidth
 - Memory size
 - I/O
 - Communication
- Low level libraries (cuDNN, MKL-DNN) have optimized memory access patterns, meaning you get the most out of your memory bandwidth
 - Make sure your DL framework is build against the appropriate low level libraries

Mitigate hardware bottlenecks

- Compute (floating point operations per second, FLOPS)
 - Memory bandwidth
 - Memory size
 - I/O
 - Communication
- Choose different architecture. E.g. TPUs, or CPUs (distributed CPU training can provide as many FLOPS as serial training on GPUs, but much more memory!)
 - Prune your model, if you can
 - Model parallelism

Mitigate hardware bottlenecks

- Compute (floating point operations per second, FLOPS)
- Memory bandwidth
- Memory size
- I/O
- Communication

Depends on what's limiting:

- IOPS: pack small files into large files. Many frameworks provide dedicated file formats for this (TensorFlows: TFRecords, Caffe: LMDB), though other packed data formats such as HDF provide similar performance benefits. Particularly important if reading from Lustre filesystem: metadataservers don't scale!
- Bandwidth: exploit parallel filesystems. E.g. use Lustre striping. See PRACE Parallel I/O best practice guide <http://www.prace-ri.eu/best-practice-guide-parallel-i-o/>
- Stage on local disks, if nodes have it. This means read from shared filesystem only occurs once – further reads (each epoch) are done from local filesystem
- Stage in RAM (/dev/shm). Only an option for small datasets.

Mitigate hardware bottlenecks

- Compute (floating point operations per second, FLOPS)
- Memory bandwidth
- Memory size
- I/O
- Communication

Depends on limiting factor:

- Latency: bundle small messages into larger ones (e.g. tensor fusion)
- Bandwidth: send as little data as possible (efficient reduction operation) using as much of the network between nodes as possible (parameter server = single bottleneck)
- Load imbalance: use homogeneous node types, all with same hardware.

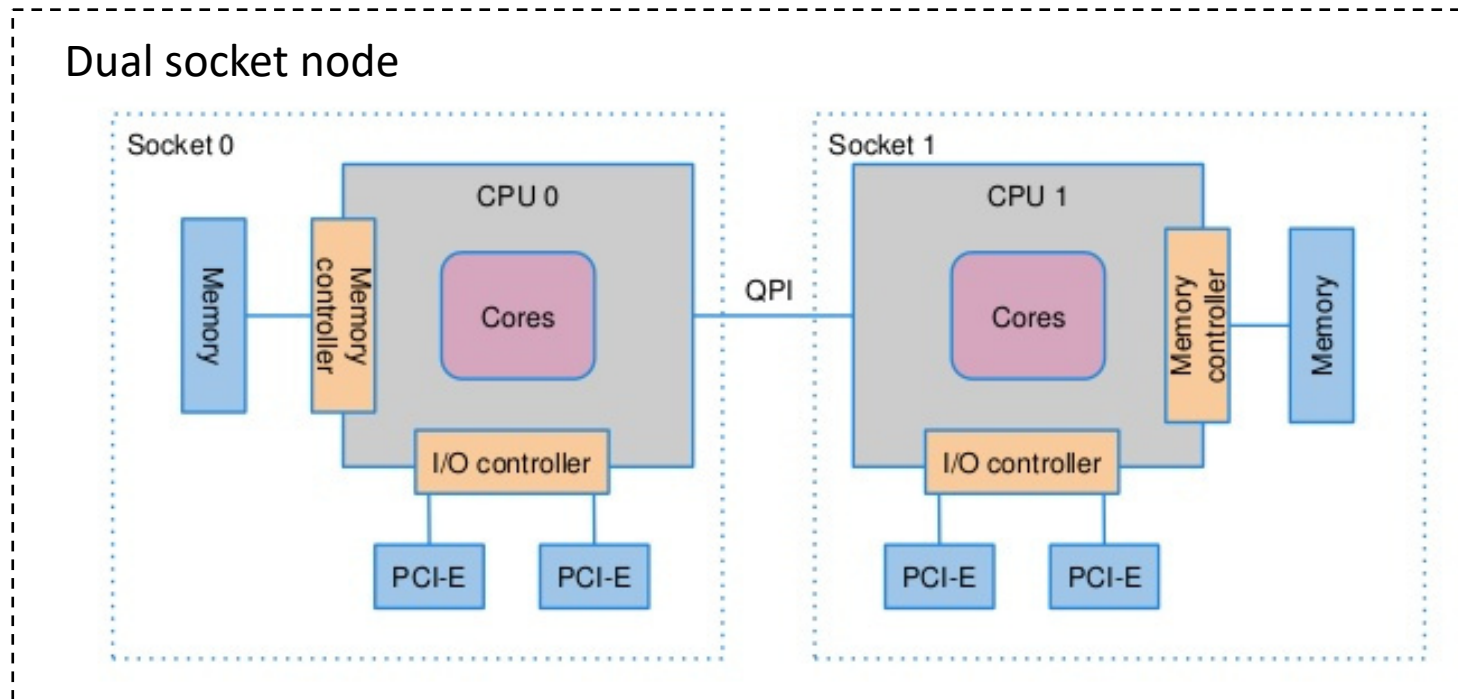
More about tensor fusion:

<https://github.com/horovod/horovod/blob/6f400014b8cb45aa013077aad0060032a4dda713/docs/tensor-fusion.rst>

Launching parallel workloads on CPU

NUMA domains

- Memory from CPU 0 and CPU 1 is 'one memory' from programmers point of view
- Cores from CPU 0 can access memory from CPU 1, but is slower!



Launching parallel workloads on CPU

Problem:

- Performance hit if process moves from CPU 0 => CPU 1 (NUMA)
- Potential performance hit when multithreading across sockets (NUMA)
- Cache misses if thread is moved from one core to another (any multithreading application)

Solution:

- Launch 1 worker per socket
- Bind (MPI) processes to sockets
- Set number of threads to number of cores available per socket
- Bind threads to cores

Launching parallel workloads

Binds threads to
(hyperthreading)cores

Set nr of threads to nr
of cores per socket

Example, 2-socket node, 12 cores per socket

- `KMP_AFFINITY="granularity=fine,compact,1,0" OMP_NUM_THREADS=12 mpirun -np 2 --map-by ppr:1:socket --bind-to socket python train.py`

Maps one process to
each socket

Binds to socket

Launches two processes

- Tip: specify ‘`--report-bindings`’ to get verbose output from *mpirun* on how processes are mapped / bound.
- Tip: specify ‘`KMP_AFFINITY="granularity=fine,verbose,compact,1,0"`’ to get verbose output on how threads are bound

Launching parallel workloads on GPU

Depends on the code!

Code designed for multi-GPU node (e.g. uses `tf.device('/gpu:0')`, `tf.device('/gpu:1')` etc):

- Launch a single process per node

Code designed for single-GPU, but parallelized with e.g. Horovod:

- Typically: launch one (MPI) process per GPU

Recap

Goals:

- Understand what hardware bottlenecks could be limiting
- Understand pro's and con's of various hardware
- Know how to choose appropriate hardware for you DL task
- Know what to do to mitigate bottlenecks

Next up:

- Learn what is the bottleneck in *your* code (profiling!)